

# Networking Architecture for Real-time Multiplayer Games

Sami Kukkonen  
University of Helsinki

## 1 INTRODUCTION

In this assignment, the task is to design a networking architecture for a multiplayer game that is played in real time. The architecture will need to be able to scale up in case the number of participants grows, reflect the actions of other users quickly and allow the user to download new game data of arbitrary sizes (*patches*). It is assumed that direct peer-to-peer connections is not always possible as players might have firewall or NAT enabled, which would block ingressing. The solution presented here does not make other assumptions about the game mechanics and is specifically focused on the networking protocol.

## 2 NETWORK TOPOLOGY

The network architecture of online multiplayer games can usually be classified into two main topologies: client-server or peer-to-peer. In client-server architectures, all players (*client nodes*) connect to a *server node* that processes events sent by clients and broadcasts updated game state to players. In peer-to-peer architecture, nodes communicate directly with each other. However, peer-to-peer architectures pose scalability issues as the clients in the game session need to be interconnected: when clients are not in the same LAN, a mechanism for discovering other nodes would be needed. In addition, as each node has to exchange traffic between all other nodes, the amount of traffic grows quadratically [1] as opposed to linearly in the client-server architecture. Figures 1 and 2 provide visualizations for these two topologies.

Using a central server instead of a peer-to-peer architecture introduces a single point of failure to the system; as the server is responsible for receiving and broadcasting updates, the capacity of the system is the same as the capacity of the server. This can be mitigated with either vertical or horizontal scaling, e.g. by load balancing packets to multiple servers and synchronizing state between them.

From a scalability perspective, a peer-to-peer architecture is clearly inferior compared to a client-server model. In addition, one of the requirements is the ability for the client to download new game configurations; a server is needed to seed the new data. Given these constraints, we choose to use a client-server topology.

## 3 TRANSPORT MECHANISM

To be able to support game sessions over the Internet, the Internet Protocol suite is chosen as the bottom level of our networking stack. This gives us two transport protocols to choose from: TCP and UDP. While TCP is best suited for transmissions where the integrity of data is important, acknowledgement and retransmission of packets might cause significant lag for the client. This is undesirable for communicating user actions or state updates because they are temporal in nature: an action such as movement or shooting is dependent on the time in which it is performed. On the other hand,




Figure 1: Peer-to-peer network topology




Figure 2: Client-server network topology

some kind of error recovery is needed for transmitting new game configurations to clients.

In order to support both temporal and perpetual data streams, we use separate UDP and TCP streams to transmit information. Temporal state updates will be transmitted via UDP and patches and initial bootstrapping will be transmitted via TCP.

### 3.1 Considerations on IP multicasting

For UDP traffic, it would be possible to use multicasting (routing the same packet to multiple destinations) instead of unicasting to reduce network traffic on the server. However, this requires substantial infrastructural support and seems infeasible in practice when communication happens over the Internet [2]. Therefore the server will use unicasting to communicate with each client.

## 4 ERROR DETECTION AND RECOVERY

UDP adds checksums to sent datagrams, which is used to ensure that the packet integrity is not compromised. However, the use of checksums is optional for IPv4 packets and required only for IPv6. [3, 4] To ensure that IPv4 connections can be used with the game, we add an additional layer of integrity verification by calculating a CRC32 checksum for each payload and prepending it to the datagram.

Congestion control is another thing that we must consider when using UDP as the transport protocol. Specifically, it is not guaranteed that UDP packets are delivered at all and if they are, the order of the received packets might not be the same than when sending. As mentioned before, packet loss is acceptable for temporal data, so we do not implement any forward error correction mechanisms. If a datagram is detected to be corrupt by a checksum mismatch, it can simply be discarded. However, out-of-order packets can cause issues when a player's state is replaced by an older state arriving later to the server.

Resiliency against out-of-order deliveries can be achieved with sequence numbering. We use two separate monotonically increasing integer sequences  $i_n$  and  $j_n$  for numbering client-originating and server-originating packets, respectively. Now the receiving side can keep count of the largest received sequence number so far and discard packets with smaller sequence numbers as stale.

To combine integrity protection and loss resiliency, we simply prepend checksums and sequence numbers to the beginning of the datagram. Thus a basic unit of transmission in the game is a triple  $(i, c, d)$  where  $i$  is the sequence number,  $c$  is the checksum and  $d$  is the transmitted data.

Let us consider an example that ties up the recovery mechanisms we've introduced: suppose that the client sends three packets  $\langle (1, c_1, d_1), (2, c_2, d_2), (3, c_3, d_3) \rangle$  to the server, but the network link is very unreliable and the server receives packets  $\langle (2, c_2, d_2), (1, c_1, d_1), (3, c_3 + 1, d_3) \rangle$ . The first packet is valid, so it is processed and the sequence counter is set to 2. The second packet has a sequence number of 1 and as  $1 < 2$  it is considered stale and discarded. The third packet has an invalid checksum, so the data is assumed to be corrupted and discarded. Now the server has processed  $d_2$ , which is a correct, although slightly outdated state.

## 5 MESSAGE FLOW

In order to present the same view of the world to all players, the players need to share the same *state*. In addition, players need to be able to send actions to the server. The server can then receive the actions from clients and compute the next state. Let  $S$  be the set of the possible states of the game world and  $\Sigma$  the set of actions. The effect of an action to the game state can be modeled as a transition function  $\delta : S \times \Sigma \rightarrow S$ . Now that the server can compute the next state of the world based on client actions, it can queue received actions, compute a new state from the current one and the actions, and send the updated state to clients; this technique is used by industrial-strength game engines such as Source [5]. The server could also update its state every time an action is received, but this would only work for scenarios where the state is cheap to compute. This might not always be the case.

Figure 3 shows an example communication flow with a 60hz server tick rate and two clients  $c_1$  and  $c_2$ . The clients send actions  $a_i \in \Sigma$  to the server. The abstract operation **UpdateState** consumes queued actions and mutates the server's state with the transition function  $\delta$ . Dashed lines represent messages sent over UDP and solid lines messages sent over a single TCP stream per client. In this scenario, the message JoinGame is required to be delivered, so it is transmitted via TCP. The client must also know the initial state

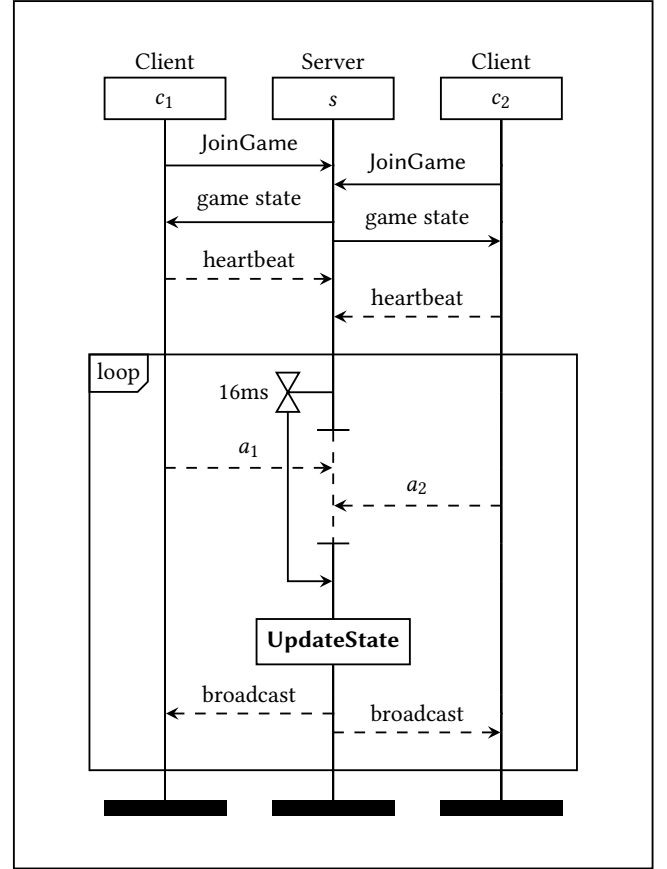


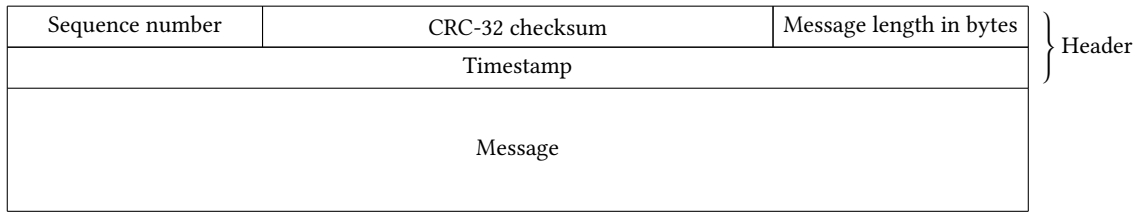
Figure 3: Sample message flow between two clients and server

of the game, so the server sends it back with TCP after the client has joined.

### 5.1 NAT and firewall considerations

Network Address Translation (NAT) is a commonly-used technique for exposing an internal network to a single public IP address via a router device. Firewalls, on the other hand, are devices or software that block connections from untrusted sources. As many routers use NAT and block ingress UDP and TCP traffic on most ports by default, they pose additional constraints to the messaging flow. In particular, it is not safe to assume that the server is able to initiate TCP connections to the client, but rather the client must initiate the TCP stream which can then be used for communication.

For UDP, the client might block datagrams sent by the server to a client if the client has not sent a datagram to the server first: on the other hand, if the client first sends a datagram to the server, NAT and firewalls will usually allow return traffic on the same port. If there is no UDP traffic from either side for an unknown period of time, the "connection" is assumed to be closed and packets are not delivered to the client. For this reason, a client must send a *heartbeat* packet to the server when joining the game and must also



**Figure 4: The structure of a single packet inside an UDP datagram**

send the heartbeat packet periodically to ensure that the firewall hole and NAT route table entry remain active.

## 5.2 Network latency

As UDP is a connectionless protocol, it is not trivial to detect whether a client has disconnected. Our architecture contains a separate TCP stream that remains open for the whole game session which can be used for evidence that the client remains connected. However, it might be useful to calculate the network latency between the client and the server. For this reason, we also include timestamps (timezone-normalized) to each sent packet. We could also use ICMP to ping both sides periodically, but combining this feature to the packets themselves seems to be a simpler solution.

## 6 CONCLUSION

The presented architecture communicates via two different protocols: TCP for persistent data for which errors have to be corrected, and UDP for temporal data for which errors have to be detected but not corrected. Using a separate TCP stream saves us the trouble of designing a custom protocol on top of UDP with the same guarantees; on the other hand, using UDP where appropriate is likely to be right choice given the real-time and latency-sensitive

nature of multiplayer games. Our application-level protocol adds ordering and integrity guarantees for datagrams and allows us to keep track of the communication latency between a client and the server. Figure 4 shows the representation of a single packet in this scheme.

## REFERENCES

- [1] Jouni Smed, Timo Kaukoranta, and Harri Hakonen. 2002. Aspects of networking in multiplayer computer games. *The Electronic Library*, 20, 87–97.
- [2] C. Diot, B.n. Levine, B. Lyles, H. Kassem, and D. Balensiefen. 2000. Deployment issues for the ip multicast service and architecture. *IEEE Network*, 14, 1, 78–88. doi: 10.1109/65.819174.
- [3] Bob Hinden and Dr. Steve E. Deering. 1998. Internet Protocol, Version 6 (IPv6) Specification. RFC 2460. (December 1998). doi: 10.17487/RFC2460. <https://rfc-editor.org/rfc/rfc2460.txt>.
- [4] Lars Eggert, Gorrry Fairhurst, and Greg Shepherd. 2017. UDP Usage Guidelines. RFC 8085. (March 2017). doi: 10.17487/RFC8085. <https://rfc-editor.org/rfc/rfc8085.txt>.
- [5] [n. d.] Source multiplayer networking. [https://developer.valvesoftware.com/wiki/Source\\_Multiplayer\\_Networking](https://developer.valvesoftware.com/wiki/Source_Multiplayer_Networking).